

Packed Qwen-Pair Row: Training Data Structure

View 1: Original text (fully visible)

"The cat sat on the warm mat by the fire"

View 2: Rewrite (densely masked input)

"A [MASK] [MASK] on a [MASK] [MASK] near the [MASK]"

Sparse focus labels:

Only a small subset of masked positions receive prediction loss



"cat"



"rested"



Dense masking removes
within-view completion clues

Key: Because the rewrite's own context is hidden, the model can only predict the sparse target words by using the visible original text — learning the source-to-rewrite correspondence.